

```
In [38]: import numpy as np
import pandas as pd
import sklearn
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [39]: ratings = pd.read_csv('data/amazon/musical_instruments_ratings.csv')
ratings.head()
```

```
Out[39]:
```

	user	item	rating
0	A2IBPI20UZIR0U	1384719342	5.0
1	A14VAT5EAX3D9S	1384719342	5.0
2	A195EZSQDW3E21	1384719342	5.0
3	A2C00NNG1ZQQG2	1384719342	5.0
4	A94QU4C90B1AX	1384719342	5.0

```
In [40]: #How many items are in the data set? How many users? How many ratings?
```

```
n_ratings = len(ratings)
n_users = ratings['user'].nunique()
n_items = ratings['item'].nunique()

print(f'Number of ratings: {n_ratings}')
print(f'Number of unique users: {n_users}')
print(f'Number of unique items: {n_items}')
```

```
Number of ratings: 10261
Number of unique users: 1429
Number of unique items: 900
```

```
In [41]: #User activity:
```

```
#What is the distribution of ratings-per-user?
```

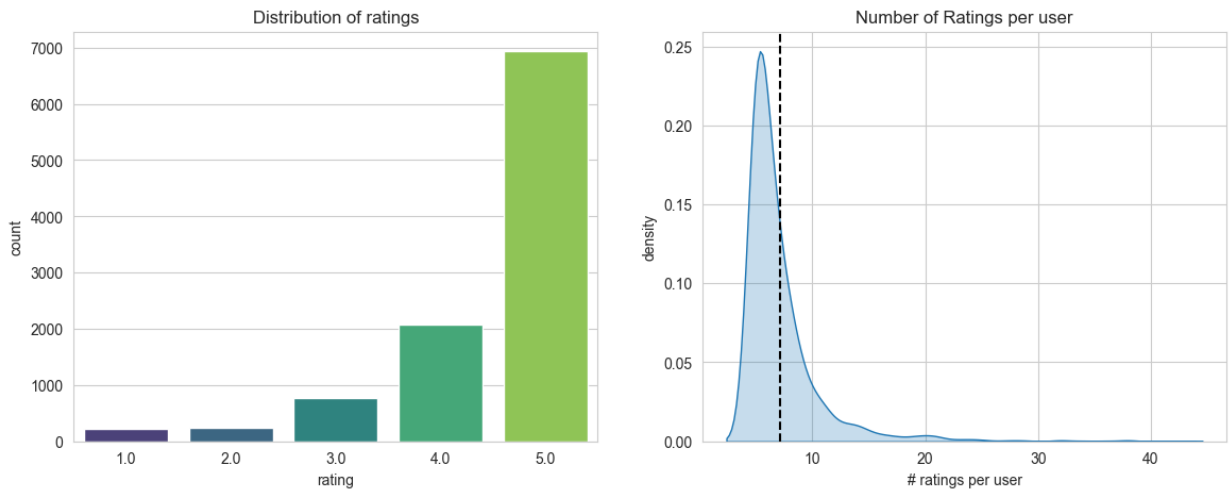
```
user_freq = ratings[['user', 'rating']].groupby('user').count().reset_index()
user_freq.columns = ['user', 'n_ratings']
user_freq.head()
```

```
Out[41]:
```

	user	n_ratings
0	A00625243BI8W1SSZNLMD	8
1	A10044ECXDUVKS	6
2	A102MU6ZC9H1N6	6
3	A109JTUXO61UY	5
4	A109ME7C09HM2M	5

```
In [42]: sns.set_style("whitegrid")
plt.figure(figsize=(14,5))
plt.subplot(1,2,1)
ax = sns.countplot(x="rating", data=ratings, palette="viridis")
plt.title("Distribution of ratings")
```

```
plt.subplot(1,2,2)
ax = sns.kdeplot(user_freq['n_ratings'], fill=True, legend=False)
plt.axvline(user_freq['n_ratings'].mean(), color="k", linestyle="--")
plt.xlabel("# ratings per user")
plt.ylabel("density")
plt.title("Number of Ratings per user")
plt.show()
```



In [48]: *#Find most and least active users for all 4 dataset. How many ratings did they*

```
user_active = ratings.groupby('user')['rating'].count()
#user_active.columns = ['n_ratings']

most_active = user_active.idxmax()
least_active = user_active.idxmin()

most_active_ratings = user_active[most_active]
least_active_ratings = user_active[least_active]

print(f"The most active user is {most_active} with {most_active_ratings} ratings")
print(f"The least active user is {least_active} with {least_active_ratings} ratings")
```

The most active user is ADH008UVJOT10 with 42 ratings.  
The least active user is A109JTUXO61UY with 5 ratings.

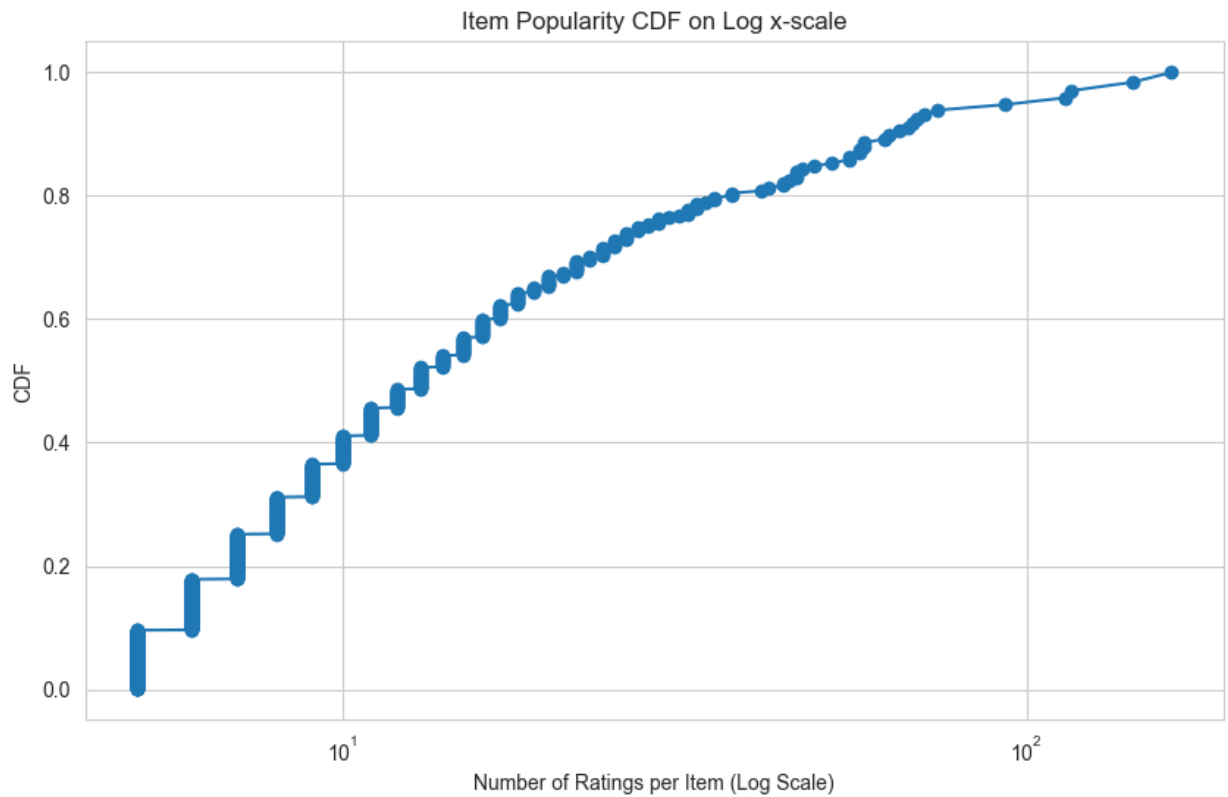
In [49]: *#Item statistics:  
#What is the item popularity curve (the distribution of ratings-per-item)? A CDF*

```
item_freq = ratings.groupby('item')['rating'].count().reset_index()
item_freq.columns = ['item', 'n_ratings']

item_freq_sort = item_freq.sort_values(by='n_ratings')

cdf = np.cumsum(item_freq_sort['n_ratings']) / np.sum(item_freq_sort['n_ratings'])


plt.figure(figsize=(10, 6))
plt.plot(item_freq_sort['n_ratings'], cdf, marker='o')
plt.xscale('log')
plt.xlabel('Number of Ratings per Item (Log Scale)')
plt.ylabel('CDF')
plt.title('Item Popularity CDF on Log x-scale')
plt.show()
```



In [50]: *#What is the distribution of average ratings for items?*

```
avg_item_rating = ratings.groupby('item')['rating'].mean().reset_index()
avg_item_rating.columns = ['item', 'avg_rating']

sns.set_style("whitegrid")
plt.figure(figsize=(10,6))
sns.histplot(avg_item_rating['avg_rating'], bins=30, kde=False, color='purple')
plt.xlabel('Average Rating per Item')
plt.ylabel('Frequency')
plt.title("Distribution of Average Ratings for Items")
plt.show()
```



```
In [52]: #Non-personalized recommendation
#What are the 10 most popular items (the items with the most ratings)? Show the

item_pop = ratings.groupby('item')['rating'].count().reset_index()
item_pop.columns = ['item', 'n_ratings']

top10_item = item_pop.sort_values(by='n_ratings', ascending=False).head(10)

print('The most popular items based on the number of ratings are: \n')
print(top10_item[['item', 'n_ratings']])
```

The most popular items based on the number of ratings are:

	item	n_ratings
706	B003VWJ2K8	163
94	B0002E1G5C	143
156	B0002F7K7Y	116
707	B003VWKPHC	114
198	B0002H0A3S	93
44	B0002CZVXM	74
258	B0006NDF8A	71
285	B0009G1E0K	69
116	B0002E2KPC	68
167	B0002GLDQM	67

```
In [53]: #What are the 10 items with the highest average ratings (with their titles and

item_avg_rating = ratings.groupby('item')['rating'].mean().reset_index()
item_avg_rating.columns = ['item', 'avg_rating']

top10_avg_rate = item_avg_rating.sort_values(by='avg_rating', ascending=False)

print('The 10 items with the highest average ratings are: \n')
print(top10_avg_rate)
```

The 10 items with the highest average ratings are:

	item	avg_rating
0	1384719342	5.0
669	B003AJMPW4	5.0
604	B0027842S4	5.0
641	B002RLLD88	5.0
288	B0009IEB0I	5.0
283	B0009EU01G	5.0
282	B0009EQOES	5.0
649	B0033P1O6S	5.0
275	B0007Y3XGW	5.0
274	B0007XWUQ2	5.0

```
In [54]: #What are the 10 movies with the highest damped average ratings, with a Bayesian
#is computed by:  $r_{ui}^{\sim} = \frac{\sum_i r_{ui}}{|R_i| + \gamma} + \gamma \bar{r}$ 

overall_mean_rating = ratings['rating'].mean()
gamma = 5

item_stats = ratings.groupby('item')['rating'].agg(['sum', 'count']).reset_index()
item_stats['damped_mean'] = (item_stats['sum'] + gamma * overall_mean_rating) / (item_stats['count'] + gamma)
item_stats['undamped_mean'] = item_stats['sum'] / item_stats['count']

top10_damped_items = item_stats.sort_values(by='damped_mean', ascending=False)

print('The 10 items with the highest damped average ratings are: \n')
print(top10_damped_items[['item', 'damped_mean', 'undamped_mean']])
```

The 10 items with the highest damped average ratings are:

	item	damped_mean	undamped_mean
594	B001W99HE8	4.897749	5.000000
504	B0015RIN6U	4.888866	4.962963
424	B000RKL8R2	4.849631	5.000000
649	B0033P1O6S	4.829581	5.000000
604	B0027842S4	4.829581	5.000000
120	B0002E2XCW	4.828874	4.866667
718	B0042EZH6W	4.826366	4.888889
387	B000L6GD04	4.820089	4.864865
438	B000RY68PA	4.817408	5.000000
669	B003AJMPW4	4.817408	5.000000

In [ ]: